

CS4186 notes

1. Image filtering

Image

- A grayscale image: a function f from \mathbb{R}^2 to \mathbb{R} , where $f(x, y)$ gives the intensity at position (x, y) .
- A digital image is a discrete (sampled and quantized) version of this function, typically stored as a matrix of intensity values (0 = black, 255 = white).
- Operators can be applied to images (transformation):

$$g(x, y) = f(x, y) + 20$$

Filters

- **Purpose of Filtering:**
 - Extract useful information (e.g., edges, contours).
 - Enhance images (e.g., remove noise, sharpen).
- **Linear Filtering (Convolution):**
 - Replace each pixel with a weighted sum of its neighbors.
 - The weights are defined by a **kernel** (or mask/filter).

Cross-Correlation and Convolution

- **Cross-Correlation ($G = H \otimes F$):**

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

- **Convolution ($G = H * F$):**

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

The kernel is flipped horizontally and vertically compared to cross-correlation.

Mean Filtering

- A simple linear filter that replaces each pixel with the average of its neighbors (moving average).
- Used for noise reduction but may blur high-frequency signals.
- example filter H :

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Gaussian Filter

- **Gaussian Kernel:**

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

The kernel weights are normalized to sum to 1.

- **Properties:**

- Removes high-frequency components (low-pass filter).
- Convolution of a Gaussian with itself yields another Gaussian.
- Less disruptive to high-frequency signals compared to mean filtering.

Sharpening

- Achieved by adding back the high-frequency details removed by blurring:

$$\text{Sharpened} = \text{Original} + \alpha \times \text{Detail}$$

where **Detail** = Original - Blurred.

Thresholding

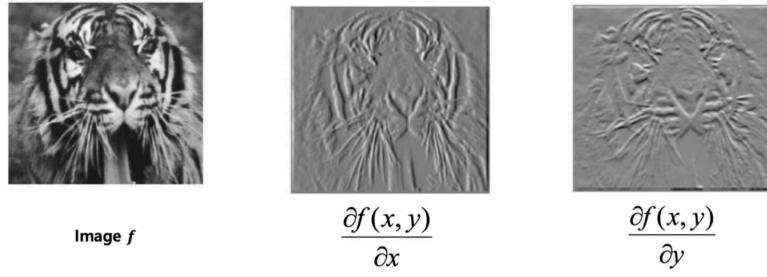
- A non-linear operation to convert grayscale images to binary:

$$g(m, n) = \begin{cases} 255 & \text{if } f(m, n) > A \\ 0 & \text{otherwise} \end{cases}$$

2. Edge Detection

- **Goal:** Identify visual changes (discontinuities) in an image, which encode semantic information.
- **Purpose:** Convert a 2D image into a set of curves for compact feature representation.

Image Derivatives



- **Finite Difference Approximation:**

- Partial derivative in x -direction:

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

- Partial derivative in y -direction:

$$\frac{\partial f}{\partial y}[x, y] \approx F[x, y + 1] - F[x, y]$$

- Implemented using linear filters (kernels H_x and H_y).

$$H_x = [-1 \quad 0 \quad 1]$$

$$H_y = \begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix}$$

Image Gradient

- **Gradient Vector:**

$$\nabla f = [\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}]$$

- Points in the direction of most rapid intensity increase.

- **Gradient Magnitude** (edge strength):

$$\|\nabla f\| = \sqrt{(\frac{\partial f}{\partial x})^2 + (\frac{\partial f}{\partial y})^2}$$

- **Gradient Direction:**

$$\theta = \tan^{-1}\left(\frac{\partial f}{\partial y}/\frac{\partial f}{\partial x}\right)$$

Smoothing and Differentiation

- **Noise Sensitivity:** Derivatives amplify noise. Solution: Smooth the image first.
- **Associative Property of Convolution**

$$\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$$

- Smoothing (h) and differentiation can be combined into a single kernel.

Sobel Operator

- Approximates the derivative of Gaussian.
- **Kernels:**

- xx direction (S_x):

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

- yy direction (S_y):

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

- Normalization (optional $\frac{1}{8}$ scaling) ensures correct gradient magnitude.

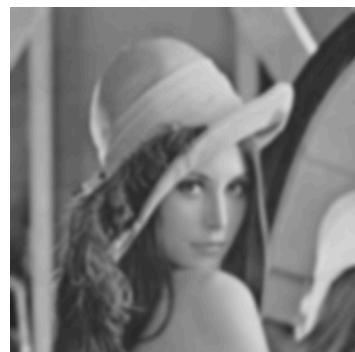
Criteria for Good Edge Detection

1. **Good Detection:** Minimize false negatives (miss real edges) and false positives (ignore noise).
2. **Good Localization:** Detected edges should be close to true edges with minimal duplication.

Canny Edge Detector

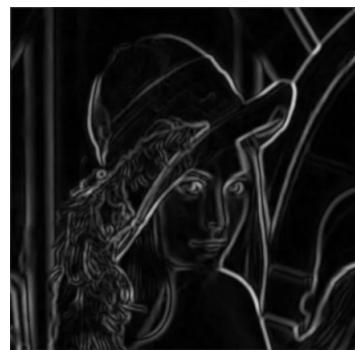


1. **Filter** image with derivative of Gaussian:

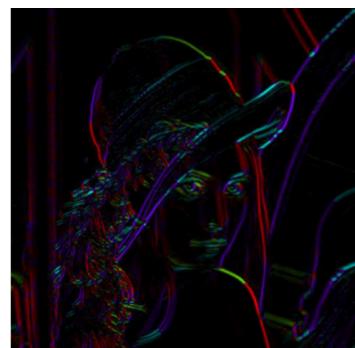


2. Compute **gradient magnitude** and **orientation**.

- **gradient magnitude:**



- **orientation:**

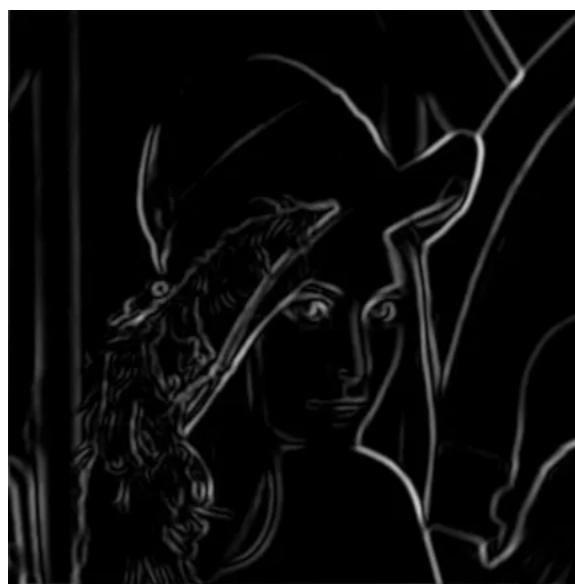


3. Apply **non-maximum suppression** to thin edges.



- Preserve local maxima in gradient direction, suppress others.

4. Perform **thresholding** (hysteresis) to link edges.



- Use high/low thresholds to distinguish strong/weak edges.

3. Image Resampling

Image Subsampling

- **Definition:** Reducing image size by discarding pixels (e.g., throwing away every other row and column).
- **Problem:** loss of high-frequency details, artifacts like moiré patterns.

Gaussian Pre-Filtering

- **Purpose:** Prevent aliasing during subsampling by smoothing the image before reducing its resolution, also help to build Gaussian Pyramid.

- **Method:**
 1. Apply a **Gaussian filter** to blur the image and remove high-frequency components.
 2. **Subsample** the smoothed image (e.g., take every 2nd pixel for half-size).

4. Color and Texture

Image Histograms (Grayscale)

- **Definition:**
 - A histogram $H[i]$ counts the number of pixels with intensity i in a grayscale image.
 - **Normalized histogram:** $P[i]$ represents the percentage of pixels with intensity i .
- **Mathematical Formulation:**

$$h(i) = \sum_x \sum_y \begin{cases} 1 & \text{if } F(x, y) = i \\ 0 & \text{otherwise} \end{cases}$$

Color Histograms

- **Purpose:** Represent the distribution of pixel values in color images for tasks like database query or classification.
- **Methods:**
 1. **3D Histogram:** Combine R, G, B channels into a single 3D histogram.
 2. **Concatenated Histograms:** Create separate histograms for each channel (R, G, B) and concatenate them.
 3. **Pseudo-Color Encoding:**
 - Encode color into a single value (e.g., 3 bits R, 3 bits G, 2 bits B \rightarrow 8-bit pseudo-color).
 4. **Normalized Color Space:** Use histograms in perceptually uniform spaces (e.g., HSV, Lab).

5. Image Texture

Edge-Based Texture Measures

1. **Edge Density:**
 - Measures "busyness" of a region by counting edge pixels.
 - Formula:

$$F_{\text{edgeness}} = \frac{|\{p \mid \text{gradient magnitude}(p) > \text{threshold}\}|}{N}$$

where N is the area size.

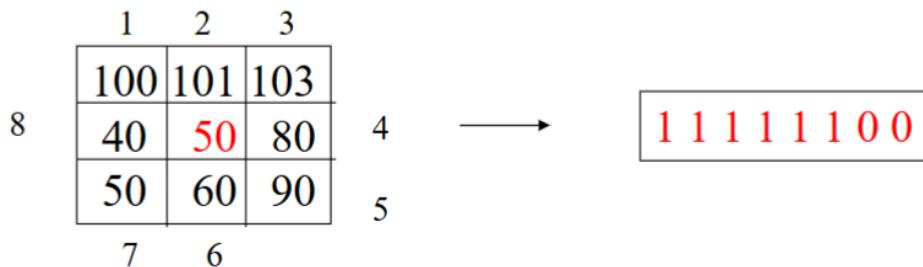
2. Edge Magnitude/Direction Histograms:

- Represents texture via histograms of gradient magnitudes ($H_{\text{magnitude}}$) and directions ($H_{\text{direction}}$).
- Feature vector:

$$F_{\text{magdir}} = (H_{\text{magnitude}}, H_{\text{direction}})$$

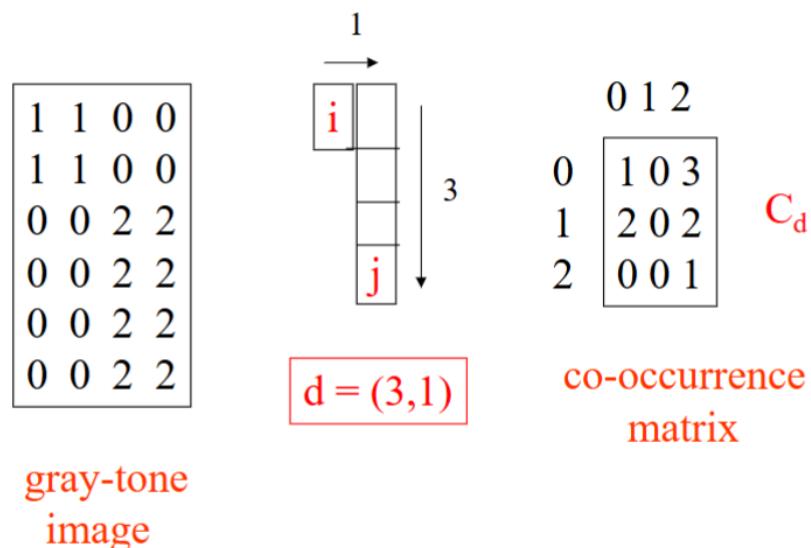
Local Binary Pattern (LBP)

- Principle:** Encodes local texture by comparing a pixel's value with its neighbors.



Co-Occurrence Matrix Features

- Definition:** A matrix $C_d(i, j)$ counts how often pixel values i and j co-occur with spatial offset $d = (dr, dc)$.



- Challenge:** Selecting optimal spatial offset d (scale-dependent).

6. Feature Extraction

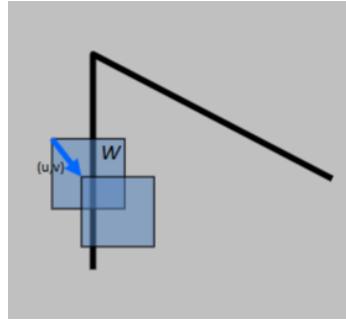
- **Purpose:** Identify distinctive points (corners, blobs) for tasks like panorama stitching, 3D reconstruction, and object recognition.
- **Key Properties of Good Features:** Locality, Distinctiveness and Invariance

Corner Detection Principle

- **Definition:** A corner is the intersection of two edges, where the image gradient changes significantly in all directions.
- **Local Uniqueness Test:**
 - **Flat region:** Minimal change in any direction.
 - **Edge:** Change along one direction (gradient perpendicular to edge).
 - **Corner:** Large change in all directions.

Harris Corner Detection

1. SSD Error for Window Shift:



- We are happy if this error is high
- For a window W shifted by (u, v) :

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

- Small motion approximation (Taylor expansion):

$$I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$$

where $I_x = \frac{\partial I}{\partial x}$, $I_y = \frac{\partial I}{\partial y}$.

2. Error Surface Approximation:

$$(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

where:

$$A = \sum_W I_x^2, \quad B = \sum_W I_x I_y, \quad C = \sum_W I_y^2.$$

- Matrix $H = \begin{bmatrix} A & B \\ B & C \end{bmatrix}$ captures local image structure.

3. Eigenvalue Analysis:

- Solve $\det(H - \lambda I) = 0$ for eigenvalues λ_1, λ_2 .

- **Corner Response:**

$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2 \quad (\text{Harris operator})$$

Alternatively:

$$f = \frac{\det(H)}{\text{trace}(H)} = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2}$$

- **Interpretation:**

- **Corner:** λ_1, λ_2 large, $\lambda_1 \sim \lambda_2, R$ large.
- **Edge:** One eigenvalue \gg the other, $R < 0$.
- **Flat:** λ_1, λ_2 small.

Harris Algorithm Steps

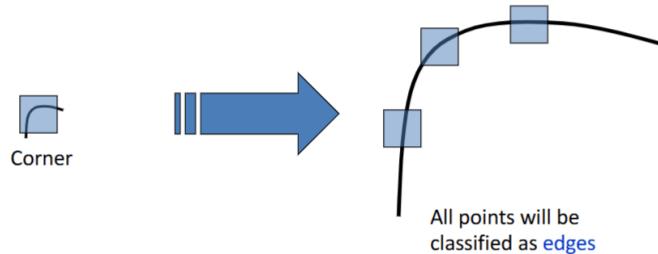
1. Compute gradients I_x, I_y (e.g., using Sobel filters).
2. Construct matrix H for each pixel:
$$H = \begin{bmatrix} \sum_W I_x^2 & \sum_W I_x I_y \\ \sum_W I_x I_y & \sum_W I_y^2 \end{bmatrix}$$
3. Compute corner response R or f .
4. Threshold R and perform non-maximum suppression to isolate corners.

Harris Detector Invariance Properties

- **Rotation Invariance**
 - Eigenvalues unaffected by image rotation.
- **Translation Invariance:**
 - Derivatives and window functions are shift-invariant.
- **Affine Intensity Invariance:**

- Invariant to intensity shift $I \rightarrow I + b$
- partially invariant to scaling $I \rightarrow aI$ (uses derivatives).

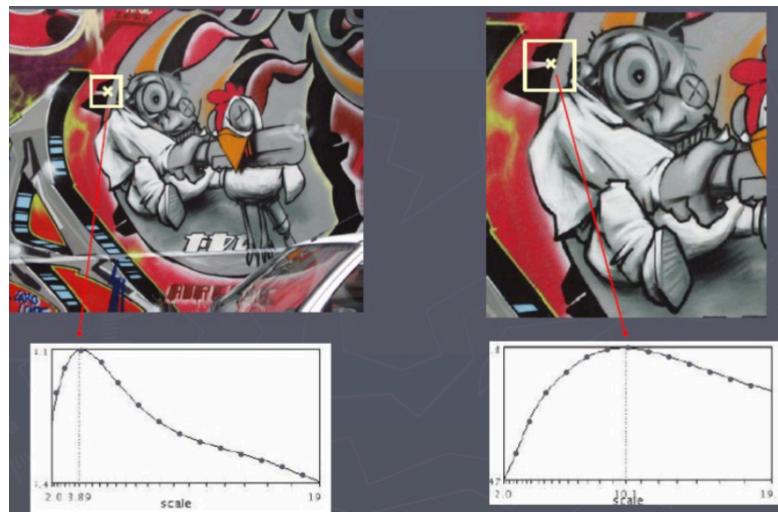
- **Scale Non-Invariance:**



- Fails under significant scaling (corners may be misclassified as edges).

Scale-Invariant Detection

- **Key Idea:** Find local maxima of a response function f (e.g., Harris operator) in **both position and scale**.
- **Automatic Scale Selection:**
 - Search for scale σ that maximizes $f(x, \sigma)$.



- **Implementation:**

- Use a **Gaussian pyramid** to handle multiple scales efficiently.

Laplacian of Gaussian (LoG)

- **Blob Detection:**
 - LoG identifies blobs as local maxima/minima in scale-space.
- **Operator Definition:**

$$\nabla^2 h_\sigma(u, v) = \frac{\partial^2 h_\sigma}{\partial x^2} + \frac{\partial^2 h_\sigma}{\partial y^2},$$

where $h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$ (Gaussian kernel).

- **Scale-Space Extrema:**

- Detect blobs by finding peaks in $\nabla^2 h_\sigma * I$ across scales σ .

Feature Descriptors

- **Purpose:** Match detected points between images.

- **Approaches:**

- **Window Matching:** Compare pixel neighborhoods.

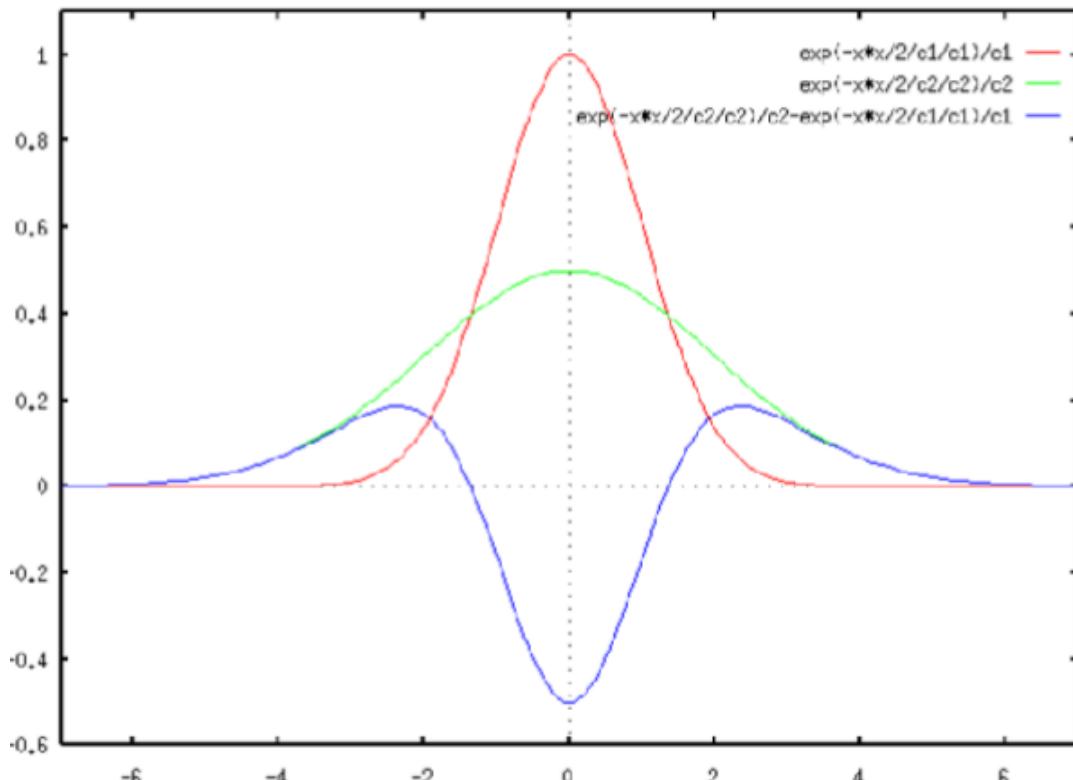
- **SIFT (Scale-Invariant Feature Transform):**

- Robust to rotation, scale, and illumination changes.

LoG-DoG (Difference of Gaussian)

- **Approximation:**

The Laplacian of Gaussian (LoG) can be approximated by the **Difference of Gaussians (DoG)** at two scales σ_1 and σ_2 (blue below):

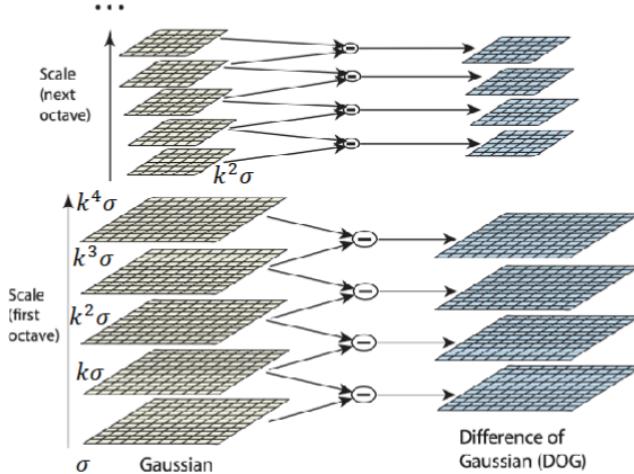


Scale Invariant Feature Transform (SIFT)

Key Steps:

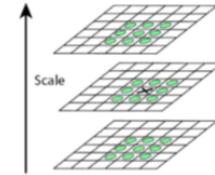
1. Scale-Space Peak Selection:

- Construct a Gaussian pyramid (multiple octaves and scales).



- Detect local maxima/minima in **DoG** space across scales and positions.

Compare a pixel (**X**) with 26 pixels in current and adjacent scales (**Green Circles**)
 Select a pixel (**X**) if larger/smaller than all 26 pixels



2. Keypoint Localization:

- Refine candidate keypoints by eliminating low-contrast or edge responses.

3. Orientation Assignment:

- Compute gradient magnitude $m(x, y)$ and orientation $\theta(x, y)$ at the keypoint scale:

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$

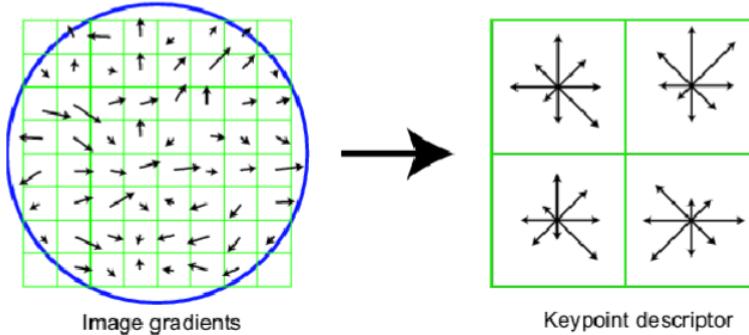
$$\theta(x, y) = \tan^{-1} \left(\frac{L(x, y+1) - L(x, y-1)}{L(x+1, y) - L(x-1, y)} \right)$$

- Assign dominant orientation(s) using a **36-bin histogram** weighted by gradient magnitude.

4. Descriptor Generation:

- Divide a **16×16** window around the keypoint into **4×4 cells**.

- Compute an **8-bin orientation histogram** per cell (relative to keypoint orientation).
- Concatenate histograms into a **128-dimensional vector** (normalized for illumination invariance).



Feature Matching

- **Distance Metrics:**
 - **SSD (Sum of Squared Differences):**

$$d(f_1, f_2) = \|f_1 - f_2\|^2$$

- **Ratio Test:**

$$\text{ratio} = \frac{\|f - f_1\|}{\|f - f_2\|}$$

where f_1 is the best match and f_2 is the second-best match in the target image.

- Rejects ambiguous matches (ratio > threshold).

7. Image Transformations

Image Warping vs. Filtering

- **Image Filtering:** Modifies pixel values (range transformation):

$$g(x) = h(f(x))$$

- **Image Warping:** Modifies pixel coordinates (domain transformation):

$$g(x) = f(h(x))$$

Forward vs. Inverse Warping

- **Forward Warping:**

- Maps source pixels $f(x, y)$ to target coordinates $(x', y') = T(x, y)$.
- **Inverse Warping:**
 - Samples target pixels $g(x', y')$ from source coordinates $(x, y) = T^{-1}(x', y')$

Parametric (Global) Transformations

- Represented as $p' = T(p)$, where T is a matrix.

- **Linear Transformations (2×2 Matrix):**

- Rotation:

$$R = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

- Scaling:

$$S = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix}$$

- **Limitation:** 2×2 matrices **cannot** represent translation.

Homogeneous Coordinates

- **Purpose:** Unify linear transformations and translation in 3D space.
- **Representation:**

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}, \quad \text{Conversion: } \begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

- **Translation:**

$$T = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

Affine Transformations

- **Definition:** Combines linear transformations and translation.

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & t_x \\ d & e & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

8. Image Alignment

Image Alignment Problem

- **Goal:** Compute a transformation \mathbf{T} that aligns two images (A and B) given a set of feature matches.
- **Approach:**
 1. Detect and match features between images.
 2. Estimate \mathbf{T} using matched points (e.g., translation, homography).

Simple Case: Translation Estimation

- **Displacement Calculation:**

For n matches, compute average translation (x_t, y_t) :

$$(x_t, y_t) = \left(\frac{1}{n} \sum_{i=1}^n (x'_i - x_i), \frac{1}{n} \sum_{i=1}^n (y'_i - y_i) \right)$$

- **Least Squares Formulation:**

Minimize sum of squared residuals:

$$C(x_t, y_t) = \sum_{i=1}^n [(x_i + x_t - x'_i)^2 + (y_i + y_t - y'_i)^2]$$

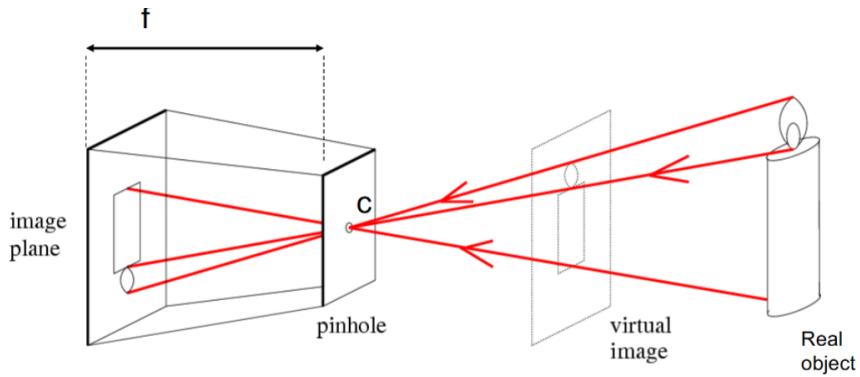
General RANSAC Algorithm (Handling Outliers)

1. **Sample:** Randomly select s points (minimal set to fit the model).
2. **Fit:** Compute model parameters (e.g., line, homography).
3. **Score:** Count inliers (points within error tolerance).
4. **Repeat:** N times, keeping the best model.
5. **Optimize:** Final least squares fit on all inliers.

9. Camera

Pinhole Camera Model

- **Principle:** Light rays pass through a small aperture (pinhole) to form an inverted image on the image plane.



f = Focal length

c = Optical center of the camera

- **Projection Equations:**

For a 3D point $P = (x, y, z)$, its 2D projection $P' = (x', y')$ is:

$$x' = f \frac{x}{z}, \quad y' = f \frac{y}{z}$$

where f is the focal length (distance from pinhole to image plane).

Homogeneous Coordinates

- **Purpose:** Simplify transformations using matrix operations.

- **Conversion:**

- 3D point:

$$(x, y, z) \Rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

- 2D point:

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- **Normalization:**

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow \left(\frac{x}{w}, \frac{y}{w} \right)$$

Camera Coordinate Systems

1. **World Coordinates P_w :** 3D scene points (x_w, y_w, z_w) .

2. **Camera Coordinates** P_c : 3D points (x, y, z) relative to the camera's optical center.
3. **Image Coordinates** P' : 2D points (x', y') on the image plane.

Intrinsic Parameters

- **Camera Matrix K** : Maps camera coordinates to pixel coordinates.

$$K = \begin{bmatrix} \alpha & 0 & c_x \\ 0 & \beta & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

where:

- $\alpha = fk$ (horizontal scale factor, k : pixels per unit length).
- $\beta = fl$ (vertical scale factor, l : pixels per unit length).
- (c_x, c_y) : Principal point (image center).

- **Projection with Intrinsics**:

$$P' = MP_c = K \begin{bmatrix} I & 0 \end{bmatrix} P_c = \begin{bmatrix} \alpha & 0 & c_x & 0 \\ 0 & \beta & c_y & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \left(\alpha \frac{x}{z} + c_x, \beta \frac{y}{z} + c_y \right)$$

Extrinsic Parameters

- **Transformation**: Converts world coordinates to camera coordinates.

$$P_c = \begin{bmatrix} R & T \\ 0 & 1 \end{bmatrix} P_w$$

where:

- R : 3×3 rotation matrix.
- T : 3×1 translation vector.

Projection Matrix (Combined Intrinsic + Extrinsic)

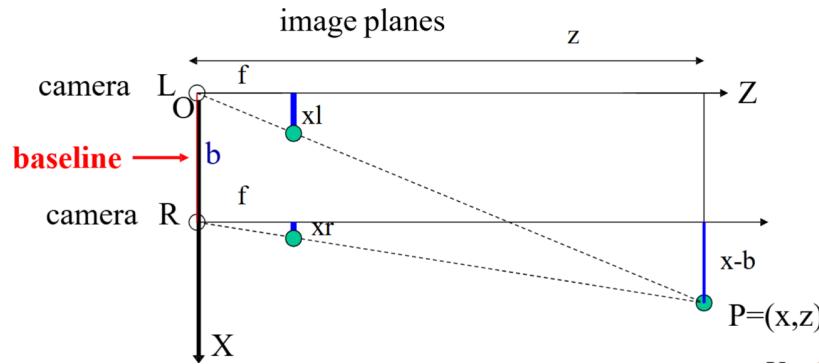
- **Full Projection**:

$$P' = K [R \ T] P_w$$

10. Stereo Vision

Depth from Stereo

- **Principle:** Disparity (horizontal shift) between corresponding points in two images is inversely proportional to depth.



- **Depth Formula:**

$$z = f \cdot \frac{b}{d}$$

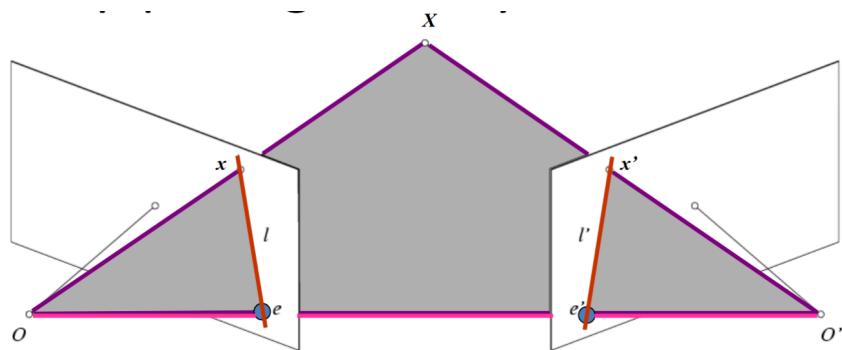
where:

- z : Depth of the 3D point.
- f : Focal length.
- b : Baseline (distance between cameras).
- $d = x_l - x_r$: Disparity (pixel difference).

- **3D Point Reconstruction:**

$$x = \frac{x_l \cdot z}{f}, \quad y = \frac{y_l \cdot z}{f}$$

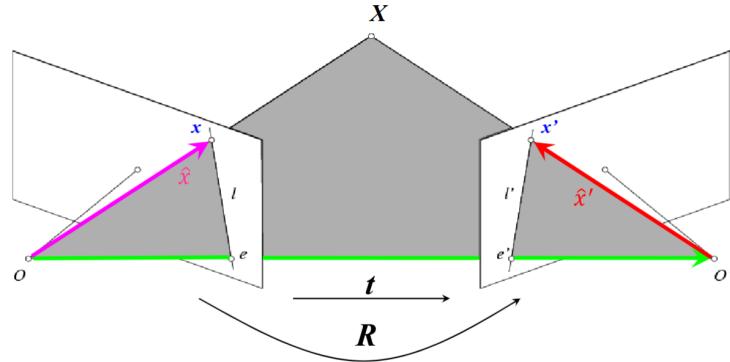
Epipolar Geometry



- **Key Terms:**

- **Baseline:** Line connecting camera centers (pink line).
- **Epipoles:** Intersections of baseline with image planes (e & e').

- **Epipolar Plane:** Plane containing baseline and a 3D point (grey area).
- **Epipolar Lines:** Intersections of epipolar plane with image planes (l & l').
- **Epipolar Constraint:** Corresponding points must lie on conjugate epipolar lines.
- **Calibrated Case:**

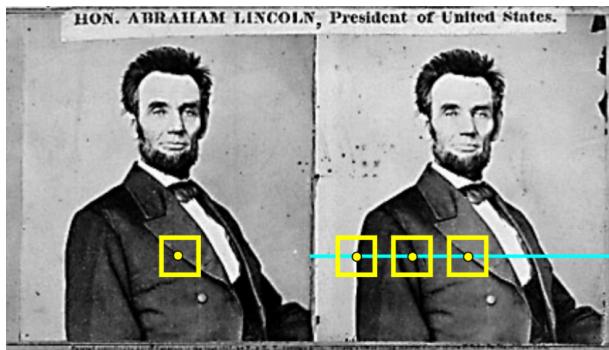


- **Assumptions:**
 - Camera intrinsics (K) and extrinsics (R, t) are known.
 - World coordinates align with the first camera's frame.
- **Projection Matrices:**
 - Camera 1: $P_1 = [I \mid 0]$.
 - Camera 2: $P_2 = [R \mid t]$.
- **Normalized Coordinates:** Remove camera intrinsics via $\hat{x} = K^{-1}x$.
- **Essential Matrix (E)**
 - **Derivation:** For corresponding points \hat{x} (Camera 1) and \hat{x}' (Camera 2), the vectors $R\hat{x}$, t , and \hat{x}' are coplanar.
 - **Epipolar Constraint:**

$$\hat{x}'^\top E \hat{x} = 0, \quad \text{where } E = [t \times] R$$

- **Properties:**
 - $E\hat{x}$ defines the epipolar line l' in the second image.
 - E encodes the relative pose (R, t) between cameras.

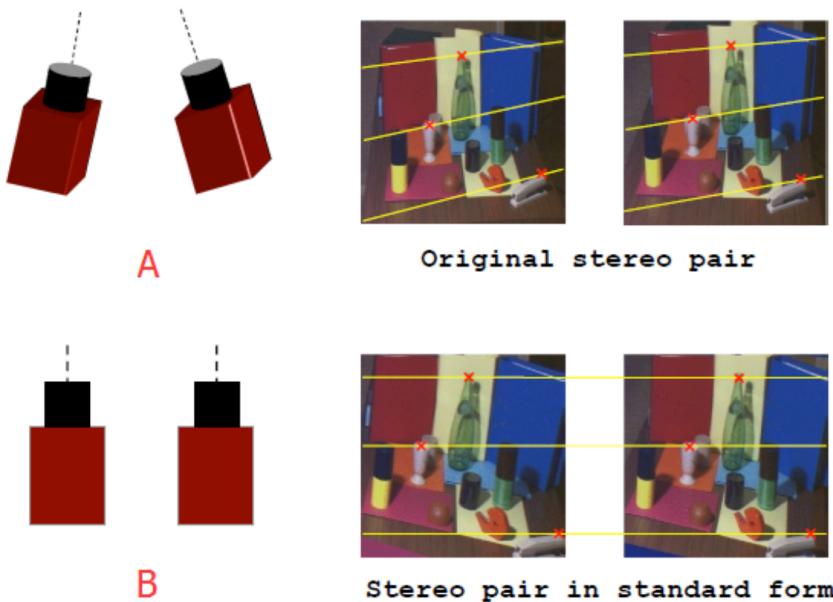
Correspondence Problem



- **Challenge:** Find matching points x (left image) and x' (right image).
- **Constraints:**
 - **Uniqueness:** Each point has at most one match.
 - **Ordering:** Matches preserve left-right order (except occlusions).
 - **Smoothness:** Disparity varies smoothly except at depth boundaries.

Stereo Rectification

- **Goal:** Align epipolar lines with image scanlines for simplified search.
- **Process:** Transform images so that corresponding points lie on the same horizontal line.



Matching Algorithms

- **Window-Based Matching:**
 - Compare patches using similarity measures:
 - **Sum of Absolute Differences (SAD):**

$$\sum_{(i,j) \in W} |I_1(i, j) - I_2(x + i, y + j)|$$

- **Sum of Squared Differences (SSD):**

$$\sum_{(i,j) \in W} (I_1(i, j) - I_2(x + i, y + j))^2$$

- **Normalized Cross-Correlation (NCC):**

$$\frac{\sum_{(i,j) \in W} I_1(i, j) \cdot I_2(x + i, y + j)}{\sqrt{\sum_{(i,j) \in W} I_1^2(i, j) \cdot \sum_{(i,j) \in W} I_2^2(x + i, y + j)}}$$

- **Trade-offs:**

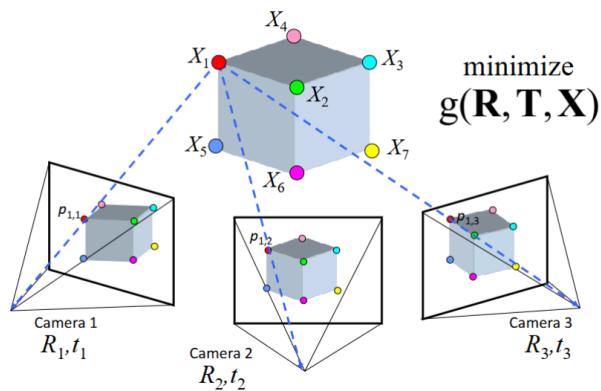
- Small windows: More detail but noisy.
- Large windows: Smoother but lose edges.

Stereo Pipeline

1. **Calibrate Cameras:** Estimate intrinsic/extrinsic parameters.
2. **Rectify Images:** Align epipolar lines.
3. **Compute Disparity:** Search for matches along scanlines.
4. **Estimate Depth:** Convert disparity to depth using $z = f \cdot \frac{b}{d}$.

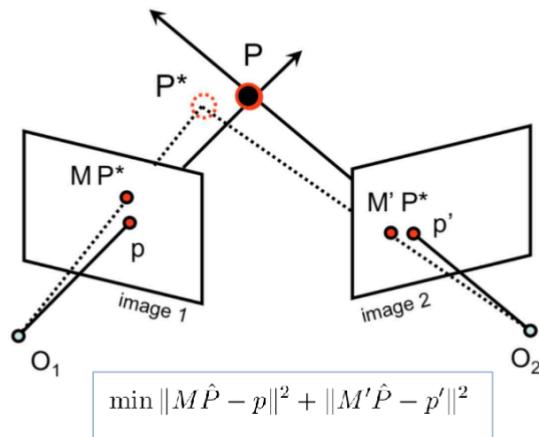
11. Structure from Motion (SfM)

- **Goal:** Reconstruct 3D scene structure and camera poses from multiple images.
- **Input:** Images with corresponding 2D feature points.
- **Objective:** Minimize reprojection error.



Reprojection Error

Re-projection Error



Key Challenges

1. Correspondence Estimation:

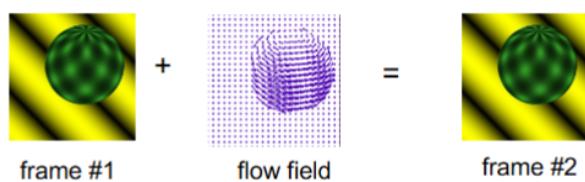
- Detect features (e.g., SIFT) and match across images.
- Use RANSAC to estimate fundamental matrices between pairs.

2. Scale of Problem:

- Large-scale SfM involves thousands of images and millions of points.
- Sparse matrices and iterative solvers (e.g., Levenberg-Marquardt) are used.

12. Optical Flow and Motion Estimation

Optical Flow Definition



- **Optical Flow:** A vector field representing the apparent motion of objects between consecutive frames, caused by relative motion between the observer and the scene.
- **Goal:** Estimate the displacement (u, v) for each pixel from $I(x, y, t)$ to $I(x, y, t + 1)$.

Key assumptions of Lucas-Kanade Tracker

- **Brightness constancy:** projection of the same point looks the same in every frame
- **Small motion:** points do not move very far

- **Spatial coherence:** points move like their neighbors

Brightness Constancy Constraint

- **Assumption:** Pixel intensity remains constant between frames:

$$I(x, y, t) = I(x + u, y + v, t + 1)$$

- **Linearized Form** (via Taylor expansion):

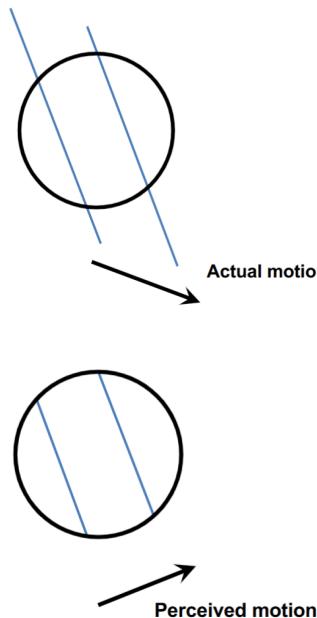
$$I_x u + I_y v + I_t = 0 \quad \text{or} \quad \nabla I \cdot [u \quad v]^T + I_t = 0$$

where:

- I_x, I_y : Spatial derivatives.
- I_t : Temporal derivative.

Aperture Problem

- **Ambiguity:** Motion along edges is underconstrained (only normal flow can be estimated).



- **Solution:** Use spatial coherence (Lucas-Kanade method) or multi-scale approaches.

Lucas-Kanade Method

Assume the pixel's neighbors have the same (u,v)

- If we use a 5x5 window, that gives us 25 equations per pixel

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix}$$

- Overconstrained linear system

$$\begin{bmatrix} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{bmatrix} \quad A_{25 \times 2} \quad d_{2 \times 1} \quad b_{25 \times 1}$$

Least squares solution for d given by $(A^T A)^{-1} A^T b$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$$A^T A \quad A^T b$$

Challenges and Solutions

1. Large Motion:

- **Coarse-to-Fine Estimation:** Use Gaussian pyramids to handle large displacements.
 - Estimate flow at low resolution, then refine at higher resolutions.

2. Iterative Refinement:

- Warp $I(t)$ toward $I(t+1)$ using current flow estimate.
- Recompute residuals and update flow iteratively.